

```
In [1]: # Import the modules required

import numpy as np # Linear algebra
import pandas as pd # data processing, CSV file I/O

import matplotlib as mpl
import matplotlib.pyplot as plt # data visualization
import seaborn as sns # statistical data visualization

# Import the dataset required

df = pd.read_csv('Month_Value_1.csv')

print(df.head(5))
```

	Period	Revenue	Sales_quantity	Average_cost	\
0	01.01.2015	1.601007e+07	12729.0	1257.763541	
1	01.02.2015	1.580759e+07	11636.0	1358.507000	
2	01.03.2015	2.204715e+07	15922.0	1384.697024	
3	01.04.2015	1.881458e+07	15227.0	1235.606705	
4	01.05.2015	1.402148e+07	8620.0	1626.621765	

	The_average_annual_payroll_of_the_region
0	30024676.0
1	30024676.0
2	30024676.0
3	30024676.0
4	30024676.0

```
In [7]: # Understand the data set and perform appropriate data cleaning
# (i) Check the data types
# (ii) Check for null values
# (iii) Check for outliers

df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 96 entries, 0 to 95
Data columns (total 5 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Period                                96 non-null     object
1   Revenue                              64 non-null     float64
2   Sales_quantity                       64 non-null     float64
3   Average_cost                         64 non-null     float64
4   The_average_annual_payroll_of_the_region 64 non-null     float64
dtypes: float64(4), object(1)
memory usage: 3.9+ KB
```

```
In [15]: # (i) Here, all the data types are appropriate
# (ii) There are some null values (with only one column filled and the rest being null)

df = df.dropna()
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Int64Index: 64 entries, 0 to 63
```

```
Data columns (total 5 columns):
```

#	Column	Non-Null Count	Dtype
0	Period	64 non-null	object
1	Revenue	64 non-null	float64
2	Sales_quantity	64 non-null	float64
3	Average_cost	64 non-null	float64
4	The_average_annual_payroll_of_the_region	64 non-null	float64


```
dtypes: float64(4), object(1)
```

```
memory usage: 3.0+ KB
```

```
In [16]: # Plot the graphs of different features against time to observe any trends and season
```

```
def plot_df(df, x, y, xlabel, ylabel, title = "", dpi = 100):
    plt.figure(figsize=(30, 4), dpi=dpi)
    plt.plot(x, y, color='tab:red')
    plt.gca().set(title=title, xlabel=xlabel, ylabel=ylabel)
    plt.show()
```

```
plot_df(df, x=df['Period'], y=df['Sales_quantity'], title='Sales quantities during the period')
plot_df(df, x=df['Period'], y=df['Revenue'], title='Revenue during the period', xlabel='Period')
plot_df(df, x=df['Period'], y=df['Average_cost'], title='Average cost during the period', xlabel='Period')
```



```
In [17]: # Augmented Dickey Fuller test (ADF Test)
```

```
from statsmodels.tsa.stattools import adfuller
```

```
data = df['Revenue'].values
```

```
pvalue = adfuller(data)[1]
```

```
if pvalue < 0.05:
```

```
    print("Series is stationary")
```

```
else:
```

```
    print("Series is non stationary")
```

```
Series is non stationary
```

```
In [18]: # converting non stationary data to stationary using differencing
```

```


df["diff_1"] = df["Revenue"].diff(periods = 1)
df.head(6)

# checking whether the Revenue feature is now stationary or not
pvalue = adfuller(df["diff_1"].dropna())[1]
if pvalue < 0.05:
    print("Series is stationary")
else:
    print("Series is non stationary")

# Plot the Revenue feature before and after differencing
df["diff_1"].plot(figsize=(14, 4));
df["Revenue"].plot(figsize=(14, 4));

```

Series is stationary




In [19]: # Compare the autocorrelation plots of the Revenue feature before and after stationari


```

from pandas.plotting import autocorrelation_plot

autocorrelation_plot(df['Revenue'].dropna())
plt.show()
autocorrelation_plot(df['diff_1'].dropna())
plt.show()

```





In [20]: *# Using the ARIMA Model for forecasting*

```
from statsmodels.tsa.arima.model import ARIMA
model=ARIMA(df['diff_1'].dropna(),order=(1,1,0))
model_fit=model.fit()
model_fit.summary()
```

C:\ProgramData\Anaconda3\lib\site-packages\statsmodels\tsa\base\tsa_model.py:471: ValueWarning: An unsupported index was provided and will be ignored when e.g. forecasting.

self._init_dates(dates, freq)

C:\ProgramData\Anaconda3\lib\site-packages\statsmodels\tsa\base\tsa_model.py:471: ValueWarning: An unsupported index was provided and will be ignored when e.g. forecasting.

self._init_dates(dates, freq)

C:\ProgramData\Anaconda3\lib\site-packages\statsmodels\tsa\base\tsa_model.py:471: ValueWarning: An unsupported index was provided and will be ignored when e.g. forecasting.

self._init_dates(dates, freq)

Out[20]:

SARIMAX Results						
Dep. Variable:		diff_1		No. Observations:		63
Model:		ARIMA(1, 1, 0)		Log Likelihood		-1113.880
Date:		Thu, 17 Nov 2022		AIC		2231.760
Time:		11:44:39		BIC		2236.015
Sample:		0		HQIC		2233.431
		- 63				
Covariance Type:		opg				
	coef	std err	z	P> z	[0.025	0.975]
ar.L1	-0.5437	0.124	-4.383	0.000	-0.787	-0.301
sigma2	2.389e+14	1.03e-17	2.32e+31	0.000	2.39e+14	2.39e+14
Ljung-Box (L1) (Q):		13.90	Jarque-Bera (JB):		2.62	
Prob(Q):		0.00	Prob(JB):		0.27	
Heteroskedasticity (H):		2.18	Skew:		0.16	
Prob(H) (two-sided):		0.08	Kurtosis:		2.05	

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).


[2] Covariance matrix is singular or near-singular, with condition number inf. Standard errors may be unstable.

In [22]: *# Using the SARIMAX Model that is a seasonal variant of ARIMA to forecast Revenue feat*

```
import statsmodels.api as sm
model = sm.tsa.statespace.SARIMAX(df['Revenue'], order=(1, 1, 1), seasonal_order=(1,1,
results = model.fit()
train_len = int(len(df['Revenue']) * 0.8)
df['Forecast'] = results.predict(start = train_len, end=len(df['Revenue']),dynamic=True)
df[['Revenue', 'Forecast']].plot(figsize=(12,8))
```

C:\ProgramData\Anaconda3\lib\site-packages\statsmodels\tsa\statespace\sarimax.py:997: UserWarning: Non-stationary starting seasonal autoregressive Using zeros as starting parameters.
warn('Non-stationary starting seasonal autoregressive'
C:\ProgramData\Anaconda3\lib\site-packages\statsmodels\tsa\statespace\sarimax.py:1009: UserWarning: Non-invertible starting seasonal moving average Using zeros as starting parameters.
warn('Non-invertible starting seasonal moving average'

Out[22]: <AxesSubplot:>



```
In [ ]:
```